

A Computational View of Medicine

Why AI will matter most where disease becomes a search problem


Birger Moëll and Erica Hughberg
Co-Founders, Eir Space

March 17, 2026

Thesis. The most important medical use of advanced AI is not generic automation. It is high-context reasoning over difficult, patient-specific problems where better synthesis can change diagnosis, treatment choice, escalation, or trial matching. In these settings, disease is better understood as a computational burden distributed across time, modalities, and decision branches than as a static label alone.

A computational framing Diagnosis as search problem

Where more inference
can improve the next
decision.



Abstract

Artificial intelligence will not improve every part of medicine equally. Its highest-value role is emerging in disease settings where the central bottleneck is not data collection alone or workflow automation alone, but high-context reasoning over multimodal evidence under uncertainty. We argue for a computational view of medicine: one that treats many serious illnesses as search problems over longitudinal records, pathology, imaging, genomics, patient-generated data, prior treatment history, and branching therapeutic options. Under this view, the key design question is not whether a model can produce clinically fluent text, but whether additional inference changes the next real decision for a specific patient. This computational constraint applies to human clinicians and AI systems alike: both face the same underlying search problem, and the quality of the search depends on the structure and depth of reasoning applied. We synthesize evidence from precision oncology, rare-disease genomics, antimicrobial discovery, sepsis surveillance, pharmacogenomics, psychiatric treatment selection, patient-reported outcome monitoring, open-notes adoption, and emerging N-of-1 medicine. We further argue that patient access to records is necessary but insufficient; because clinical reasoning is in practice distributed across many specialists and institutions—with the patient as the only persistent entity across all episodes—the practical value of data access rises sharply when patients and clinicians can repeatedly compute over the full timeline of illness. The result is a framework for directing AI toward cases where reasoning is the bottleneck, outcomes are sensitive to synthesis quality and calibration, and patient-held data can become a force multiplier rather than a static archive.

1 Introduction

Medicine has long been organized around disease labels, organ systems, and specialty workflows. Those abstractions are useful, but they often obscure where the real difficulty lies. In many of the cases that matter most, the limiting factor is neither the absence of a diagnosis code nor the absence of a single decisive test. It is the burden of synthesis: reconciling longitudinal evidence, deciding which signals matter, ranking plausible interventions, and updating the model of the case as new data arrive.

This paper argues that advanced AI should be judged through that lens first. The central question is not whether a system can imitate clinical language. It is whether additional inference changes a meaningful next decision in settings where the decision is difficult, high-stakes, and patient-specific. This framing is consistent with the broader case for human-plus-machine medicine made by Topol, but it places computational leverage rather than generic augmentation at the center [31].

This is not the first time medicine has undergone a computational upgrade. Evidence-based medicine, which reshaped clinical practice from the 1990s onward, was itself a shift in how reasoning was done: it replaced individual recall and apprenticeship-derived judgment with structured synthesis over aggregated trial data. Cochrane reviews are, in effect, pre-computed answers to clinical questions. The computational view proposed here is the next step in the same progression. Where EBM asked “what does the average patient in this trial tell us about this patient?”, the computational view asks “what does the full multimodal, longitudinal record of *this* patient tell us about *this* decision?” Each transition has expanded the data surface and the reasoning capacity available to the clinician. The question is no longer whether medicine should be computational, but how deep the computation should go and where to direct it.

Crucially, this computational constraint is not unique to machines. It is already the daily reality of clinical practice. Physicians employ both rapid pattern recognition and slower deliberate analysis—what cognitive science calls dual-process reasoning—and they switch between these modes depending on case complexity [8]. But the time a clinician can spend thinking about any single patient is finite. If a doctor had twice as long to reason through a complex case, would they reach a better conclusion? In many routine situations, the answer is no: the protocol is clear and more thought adds nothing. But in high-burden cases—a confusing constellation of symptoms, an ambiguous genomic result, a cancer that has stopped responding—more

reasoning time genuinely changes the output. This is not merely a theoretical claim. The “diagnostic timeout” in emergency medicine—a deliberate, structured pause to reconsider the differential diagnosis—has been shown to reduce diagnostic error, providing direct empirical evidence that allocating more cognitive compute to a case works in practice [8, 10]. Furthermore, primary care consultation lengths vary sharply across international healthcare systems, ranging from 5 minutes in high-volume settings to more than 20 minutes in others [15]. If reasoning time is the central bottleneck the paper claims it is, this variation should predict diagnostic quality, grounding the computational metaphor directly in healthcare economics.

Recent work on reasoning-focused language models suggests that AI systems can engage in structured clinical reasoning that parallels this deliberate process, including hypothesis generation, differential weighting, and evidence integration. However, the relationship between reasoning length and decision quality is not monotonic. Evaluations of models like DeepSeek R1 reveal a clear tension: shorter reasoning chains often correlate with accuracy, while excessively long chains correlate with errors driven by anchoring bias, overthinking, and rationalizing wrong conclusions [22]. The value of additional reasoning depends heavily on its functional form. It is productive when it widens the search to explore new hypothesis space, but it becomes destructive when it merely deepens commitment to a single, incorrect branch. This distinction separates good clinical reasoning from bad clinical reasoning in humans, and it turns “more reasoning helps” from a slogan into a concrete design principle for medical AI. The implication is not that machines reason exactly as clinicians do, but that diagnosis and treatment selection are search problems, and the quality of the search depends on the structure, not just the duration, of the reasoning applied.

The argument has three parts. First, disease classes differ sharply in how much more computation can improve outcomes. Second, patient access to longitudinal records changes the practical location of medical intelligence: once records can be assembled across sites, compute can move with the patient rather than remaining trapped inside an institution. Third, the future of high-value health AI is likely to look increasingly N-of-1: iterative, longitudinal, individualized, and explicitly tied to decisions rather than one-shot answers.

Definition. A *computational view of medicine* treats disease management as the problem of converting fragmented, multimodal, longitudinal evidence into better patient-specific decisions under uncertainty.

2 Scope and evidence selection

This paper is a scientific perspective rather than a systematic review. The evidence base was assembled to support a specific claim: that the clinical value of AI is concentrated in disease settings where the main bottleneck is patient-specific synthesis under uncertainty. We therefore prioritized five categories of evidence: prospective interventional studies in which computational systems changed care operations or outcomes; landmark mechanistic or translational studies that make the search-space argument concrete; studies on patient access to records and patient-generated data; literature on N-of-1 and individualized therapeutic design; and implementation-oriented work on digital phenotyping, quality, and safety [14, 17, 20].

The paper does not claim that all cited domains already have mature evidence for the full computational paradigm. Instead, it distinguishes three layers of support: direct interventional evidence that computation changed outcomes; translational evidence that a disease domain is intrinsically computational; and forward-looking inference about how patient-held data and repeated synthesis could alter future care models. Where the paper makes a prospective claim rather than reporting demonstrated benefit, that claim should be read as an explicit extrapolation from current evidence rather than as settled fact.

3 A framework for computational leverage

We propose five properties that determine whether a disease class is likely to benefit from larger inference budgets:

1. **Search-space size:** how many plausible hypotheses, targets, interventions, or sequences must be evaluated.
2. **Data burden:** how many modalities and how much longitudinal context are required to reason well.
3. **Personalization:** how much the correct answer depends on the specific patient rather than the average protocol.
4. **Actionability:** whether better synthesis can materially alter diagnosis, escalation, treatment choice, or trial eligibility.
5. **Compute elasticity:** whether additional reasoning depth is likely to improve the result instead of merely restating the same answer more fluently.

Under this framework, precision oncology, rare-disease genomics, resistant infection, certain critical care syndromes, and selected long-horizon chronic conditions rise to the top. These domains combine large search spaces with heterogeneous data and decisions whose quality directly affects outcome. By contrast, diseases governed primarily by narrow protocols or limited action space may benefit more from automation than from massive inference.

This framework also clarifies when AI adds value beyond what a clinician can already provide. In domains with low compute elasticity, a clinician working at normal pace already reaches the correct decision—more thinking time, whether human or machine, buys nothing. In high-elasticity domains, every additional minute of deliberation can shift the differential, reorder treatment options, or surface a trial that would otherwise be missed. The practical reality is that clinicians cannot spend an hour reasoning through every case, even when an hour would help. AI does not replace clinical reasoning in these settings. It extends the reasoning budget available for the cases that need it most.

To appreciate why compute elasticity matters so much, consider the sheer magnitude of the reasoning gap in current practice. A molecular tumor board typically spends 15–30 minutes per case. The relevant decision space includes hundreds of approved oncology drugs, thousands of active clinical trials, and millions of published papers; the team searches perhaps a fraction of one percent of that space in the time available. A geneticist reviewing a rare disease case for the first time may have 60 minutes and a patient file accumulated over years of prior workup; the relevant variant literature alone runs to millions of entries. A primary care physician managing a complex multimorbid patient on twelve medications has 15 minutes; the combinatorial space of drug–drug, drug–disease, and disease–disease interactions is enormous, and the doctor navigates it almost entirely from memorized heuristics rather than computed analysis. These are not speculative figures. They represent a measurable *compute deficit*: the gap between what is searchable and what is actually searched. The computational view argues that closing this gap—not automating paperwork—is where the largest clinical returns will come from.

Crucially, while the scores in Table 1 are conceptual illustrations, they point toward a concrete, falsifiable research programme. These dimensions can be operationalized empirically: *search-space size* can be measured by counting the number of plausible differentials, approved therapies, or the relevant variant space for a typical case; *compute elasticity* can be quantified through experimental designs that systematically vary reasoning time or language-model inference depth, measuring the resulting change in downstream

Domain	Search	Data	Pers.	Action	Elast.	Interpretation
Precision oncology	5	5	5	5	5	Dense multimodal evidence and many branching intervention paths make more inference directly decision-relevant.
Rare-disease genomics	5	5	5	4	5	Large hypothesis spaces and sparse distributed clues create unusually high value for long-context reasoning.
Resistant infection	4	4	4	5	4	The answer depends on organism, host state, prior exposure, and time-sensitive escalation.
Sepsis / ICU	4	5	4	5	4	Continuous streaming data create a high-burden monitoring problem in which timing dominates outcome.
Chronic inflammatory disease	3	4	4	4	3	Longitudinal symptom and biomarker integration can improve timing and sequencing, but action spaces are narrower.
Routine protocol care	1	2	1	2	1	Most value lies in reliable execution and logistics rather than larger inference budgets.

Table 1: Illustrative scoring matrix for computational leverage. Scores are conceptual rather than empirical, but they make the framework explicit and falsifiable.

decision quality; and *data burden* can be assessed by counting the distinct data modalities routinely required before a defensible clinical decision can be made. Defining these metrics transforms the computational view from a descriptive metaphor into an actionable engineering target.

Why prevalence is not enough. Common diseases are not automatically the best targets for large inference budgets. The relevant question is whether more synthesis can change a meaningful patient-specific decision. That is why prevalence rankings and compute-opportunity rankings diverge.

4 Why patient-held data matter

The computational view becomes much stronger when we recognize that centralized medicine is already a distributed computation. The paper thus far largely frames reasoning as a single-agent problem—"a clinician" reasoning about a case. In reality, clinical reasoning is distributed across general practitioners, specialists, pathologists, radiologists, and multidisciplinary team meetings. Each referral is a computational step, and each handoff is a potential information loss. The "diagnostic odyssey" in rare disease—averaging 7 years and 15 referrals to reach a diagnosis—is not a story about a single doctor who failed to think hard enough. It is the story of a distributed search that could not maintain state across its own nodes. Hypotheses get dropped at specialty boundaries, and prior exclusions are repeated because no single entity holds the full search history.

This reframing reveals what AI is uniquely positioned to do: not merely "think longer" in isolation, but hold the full search state together across the entire trajectory of fragmented care. This connects directly to the




Figure 1: The patient-compute loop. In high-burden disease, value comes from repeated synthesis over the full timeline rather than one-shot question answering.

necessity of patient data access. The patient is the only entity that persists across all these episodes, making them the natural locus of computational state. Once the patient can assemble data across encounters, institutions, and modalities, compute can move with the patient rather than remaining trapped inside an institution. OpenNotes demonstrated that patients value direct access to clinicians' notes and that transparency can improve engagement with care [9]. Follow-on work showed that patients often shared those notes with family members and caregivers, turning the note into a coordination artifact rather than a static document [16]. In parallel, patient-generated data streams have become increasingly common across chronic disease, symptom tracking, wearables, and home monitoring [24]. The strategic implication is that records are no longer only institutional memory; they are becoming a substrate for patient-held, continuous reasoning.

Raw access, however, is not enough. Tuttle and colleagues argued that precision medicine will remain incomplete if patients are not active participants in how data are interpreted and used [32]. Hassan and colleagues similarly emphasize that AI for shared decision-making must improve understanding and deliberation, not simply generate recommendations [12]. Jim and colleagues reviewed the analytic and operational landscape of patient-generated health data in oncology and concluded that the integration challenge is not merely technical but also clinical, organizational, and evidentiary [17]. More recently, Salmi and colleagues showed in a proof-of-concept study that patients can already use large language models on open notes to interpret and work with their records [28]. That result should be read carefully: it is early evidence, not mature clinical validation. But it clarifies the direction of travel. Once patients can access timelines, test results, notes, and symptom logs, the missing layer is repeatable reasoning over that history.

4.1 Counterarguments and constraints

Patient-held data are not automatically clinically useful. Provenance may be uncertain, timestamps may be incomplete, self-tracked measures may drift in quality, and the patients most likely to benefit from longitudinal computation may also face the largest barriers in digital access. Huckvale and colleagues make a similar point for digital phenotyping: clinically interesting data streams do not become clinically usable unless quality, safety, equity, and implementation needs are addressed directly [14]. The computational paradigm therefore depends on a stricter standard than mere interoperability. It requires provenance, chronology, validation, and interfaces that separate observed evidence from inferred conclusions.

Access versus leverage. Data access creates possibility. Computational interpretation creates leverage. In a patient-centered health stack, the two should be designed together.

5 Evidence from the highest-leverage domains

5.1 Precision oncology

Cancer is already a computational problem. Molecular diagnosis, resistance interpretation, treatment sequencing, and trial matching require synthesis across pathology, imaging, genomics, prior therapy, and time. Personalized cancer vaccines make that explicit: the workflow depends on identifying patient-specific tumor mutations, ranking neoantigens, and manufacturing a bespoke intervention. In a landmark study, Sahin and colleagues showed that individualized RNA mutanome vaccines could mobilize poly-specific immunity against melanoma [27]. Structure-prediction systems such as AlphaFold further expand the search surface by making protein consequence analysis more tractable across large molecular spaces [19].

The importance of compute here is not cosmetic. It sits inside the decision itself: which mutation matters, which escape mechanisms are plausible, which combination is worth pursuing, and which trial remains defensible for this patient now. Oncology is one of the clearest examples of a disease domain where the burden of reasoning has become part of the disease burden itself.

A concrete workflow makes the point clearer. The same patient may generate surgical pathology, radiology progression patterns, whole-exome or panel sequencing, circulating tumor DNA, treatment toxicities, and eligibility-relevant comorbidities. The computational task is to rank candidate drivers, resistance hypotheses, and therapeutic branches while making the uncertainty around each step visible. In other words, the leverage does not come from generating prose about cancer. It comes from compressing a large search problem into a smaller set of defensible next moves.

5.2 Rare disease genomics

Rare-disease diagnosis is where long-context reasoning becomes immediately clinically useful. The problem is not only finding a variant. It is aligning phenotype, pedigree, prior negative workup, and the literature fast enough to affect care. In 2024, Wojcik and colleagues reported in the *New England Journal of Medicine* that genome sequencing materially improved diagnosis for patients with suspected rare disease and supported its use as a first-tier test in appropriate settings [33].

This is a canonical compute-heavy domain: large hypothesis spaces, sparse signal, fragmented records, and a direct payoff when reasoning succeeds. The same logic applies to critically ill infants, undiagnosed neurologic disease, and multisystem pediatric cases where delay compounds. When diagnosis arrives late, the cost appears as additional procedures, avoidable admissions, and years spent navigating the wrong branch of care.

5.3 Resistant infection and antibiotic discovery

In antimicrobial resistance, the search problem has two layers. The first is patient-specific treatment selection: organism identity, susceptibility, host state, prior exposure, organ function, and drug interactions. The second is discovery itself: how to search chemical space for compounds with novel activity. Stokes and colleagues used deep learning to identify halicin, demonstrating that machine learning can materially accelerate antibiotic discovery in regions of chemical space that were not obvious to standard pipelines [30].

This is exactly the kind of problem where more compute can change outcomes: better regimen ranking in the clinic and better exploration of candidate space in the lab.

5.4 Sepsis and critical care

Critical care is different from oncology and rare disease, but it is no less computational. Here the challenge is time. Dozens of measurements update continuously, and the relevant signal is often distributed across trends rather than single observations. Adams and colleagues showed that a machine-learning sepsis early-warning system deployed prospectively across multiple hospitals was associated with faster antibiotic administration and lower mortality [2]. Rajkomar and colleagues had earlier shown that deep learning on electronic health records could scale across broad clinical prediction tasks using the raw longitudinal record itself [26]. Moses and colleagues extended this logic into pragmatic deterioration surveillance, showing that real-time systems can be attached to live escalation pathways [23].

These results matter because they establish two parts of the same claim: longitudinal records contain clinically actionable signal, and systems that can hold more of that context together can improve timing-sensitive decisions.

5.5 Psychiatry and mental health

While oncology, rare disease, and infection have relatively well-defined targets, psychiatry is arguably the most computationally complex domain in medicine. Mental health care involves enormous hypothesis spaces, longitudinal and multimodal data, and highly individual treatment responses (for example, first-line antidepressants fail in approximately 50% of cases). The current standard of care functions essentially as a brute-force search through medication options over months or years, with each failed trial costing time, side effects, and continued illness.

Pharmacogenomics offers a partial computational shortcut. Genetic variation in drug-metabolizing enzymes—most notably the cytochrome P450 family—can predict whether a patient will metabolize a given medication too quickly, too slowly, or normally. When this information is available before prescribing, it can eliminate options likely to fail and prioritize those more likely to succeed. The Clinical Pharmacogenetics Implementation Consortium has published dosing guidelines for dozens of drug–gene pairs, and evidence for genotype-guided prescribing is accumulating across psychiatry, cardiology, and oncology [13]. The bottleneck here is not data collection—genotyping is cheap and fast—but integration: connecting the genotype to the drug list to the patient’s history to the prescribing decision. This is a synthesis problem, not a sequencing problem.

But the computational burden in psychiatry extends well beyond drug selection. The deeper reason the problem is so hard is that the relevant search space is not a genome or a set of lab values. It is the patient’s entire cognitive and experiential landscape: the learned patterns, coping strategies, relational templates, and interpretive habits shaped by a lifetime of experience. Many mental health conditions can be understood as maladaptive but semi-stable states—ways of interacting with the world that, although counterproductive, persist because they are locally reinforced. A person with social anxiety avoids social situations, which reduces short-term distress but deepens long-term isolation. A person with depression withdraws from activity, which feels protective but eliminates the sources of positive reinforcement. These are local minima in a behavioral landscape: stable enough to resist casual intervention, but not globally optimal.

Cognitive behavioral therapy makes this computational structure explicit. The ABC model—antecedent, behavior, consequence—is a structured search through the patient’s cognitive patterns: identify the triggering situations, map the behavioral responses, trace the reinforcing consequences, and find the self-sustaining loops that maintain the disorder [6]. The therapist’s task is, in the vocabulary of this paper, a search problem over a high-dimensional space defined by the patient’s beliefs, behaviors, and life context. The goal is to locate the specific patterns that, if restructured, would shift the patient out of the maladaptive equilibrium and into a more adaptive one. This is precisely the kind of problem where more reasoning—more

careful exploration of the patient's cognitive landscape—changes the outcome.

The scale of latent demand for this kind of computational support is already visible. Surveys consistently report that a large fraction of users of general-purpose large language models use them for emotional support, self-reflection, and informal therapeutic conversation. This is not clinical evidence, and it should not be mistaken for validated therapy. But it reveals something important: people already intuit that working through their own cognitive patterns is a reasoning process, and they reach for computational tools when human reasoning time is unavailable or insufficient.

Although objective biomarkers are often sparser in psychiatry than in oncology, the fundamental computational burden—iterative search over the full landscape of a person's cognition, behavior, and experience—makes psychiatry not a marginal case for the computational view but one of its most demanding and important test cases.

5.6 Patient-reported outcomes and symptom monitoring

Not all high-value computation depends on genomics. Basch and colleagues demonstrated that electronic patient-reported symptom monitoring during cancer treatment improved overall survival, likely by enabling earlier clinical response to deterioration patterns that would otherwise be missed [4, 5]. This is an important bridge between institutional and patient-held data. Symptoms, side effects, and treatment tolerability often live outside the hospital until they are severe enough to trigger contact. In a computational paradigm, those signals become first-class inputs into the case model.

6 N-of-1 medicine and personal computation

The computational view converges naturally with N-of-1 medicine. Schork argued that precision medicine should increasingly be understood at the level of the individual rather than the cohort average [29]. That argument has a formal clinical-trial lineage. Guyatt and colleagues laid much of the early foundation for randomized trials in individual patients [11], and Lillie and colleagues later argued that N-of-1 trials are unusually well aligned with the actual endpoint of clinical practice: optimizing care for a single person [20]. More recent work has pushed this further. Augustine and colleagues proposed a roadmap for N-of-1 therapies, describing the scientific, operational, and regulatory requirements for highly individualized interventions [3]. Jonker and colleagues similarly outline how N-of-1 medicine could reshape the development and evaluation of treatments for rare and heterogeneous disease [18]. These papers matter because they move the discussion from personalization as a slogan to personalization as an organizing scientific unit.

Longitudinal self-measurement provides a parallel but distinct line of evidence. Chen and colleagues showed in 2012 that integrative personal omics profiling could reveal dynamic molecular and medical phenotypes within a single individual over time [7]. Piening and colleagues later demonstrated that personal, dense, dynamic data clouds can be used to monitor health transitions with much greater temporal resolution than standard episodic care [25]. These studies do not imply that every patient should become a self-quantification project. They do show that biology, symptoms, and risk are often more legible in longitudinal within-person data than in sparse snapshot encounters.

Public self-experimentation should be interpreted separately again. Highly instrumented self-tracking programmes such as Bryan Johnson's are not clinical evidence and should not be treated as such. What they do reveal is latent demand: some people already want to run dense longitudinal data loops over their own bodies. The scientific path forward is not celebrity optimization. It is the translation of that demand into audited N-of-1 workflows that combine patient goals, validated measurements, and clinically accountable

Domain	Core data	How computation is leveraged	Representative evidence
Precision oncology	Genomics, pathology, imaging, treatment history	Variant interpretation, resistance modeling, neoantigen ranking, regimen sequencing, trial matching	Sahin et al. [27]; Jumper et al. [19]
Rare disease	Phenotype timeline, pedigree, genome sequencing, prior workup	Cross-linking phenotype ontologies, literature, family history, and candidate variants	Wojcik et al. [33]
Resistant infection	Microbiology, susceptibilities, organ function, prior antibiotic exposure	Better ranking of therapies and faster search across compound or regimen space	Stokes et al. [30]
Sepsis / ICU	Streaming vitals, labs, notes, medication events	Earlier escalation from weak, time-distributed patterns across the longitudinal record	Adams et al. [2]; Rajkomar et al. [26]; Moses et al. [23]
Patient-reported symptoms	Patient-reported outcomes, treatments, toxicities, goals	Detecting deterioration and treatment intolerance before they become acute	Basch et al. [4]; Basch et al. [5]
Patient-held records	Notes, meds, labs, imaging reports, symptoms, goals	Converts raw data access into usable reasoning leverage before and between encounters	Delbanco et al. [9]; Salmi et al. [28]

Table 2: High-value disease areas under a computational lens.

reasoning.

7 Case examples

Table 2 shows why a single computational model fits these otherwise different domains.

7.1 Concrete pathways for future use

Figure 2 makes the future-oriented claim concrete. In each case, the point of computation is not to replace the clinician. It is to improve preparation, ranking, and timing for the next branch in the care pathway.

1. **Precision oncology pathway.** A patient-assembled case file combines sequencing, pathology, imaging, and prior therapies. Compute is used to rank candidate drivers, plausible resistance mechanisms, toxic trade-offs, and live trials. The output is not “the answer” but a defensible next branch with explicit uncertainty.
2. **Rare disease pathway.** Parents or clinicians maintain a phenotype timeline with milestones, regressions, imaging, and prior exclusions. Compute links that timeline to genome evidence and the literature, reducing time spent re-deriving the case from scratch at each tertiary referral.

Future patient personas		
Patient	Data substrate	Computational leverage
Asha, 58 metastatic NSCLC	Tumour DNA/RNA, pathology, CT trends, prior lines, toxicities, performance status	Variant prioritization, resistance hypotheses, neoantigen ranking, and live trial matching change the next therapeutic branch.
Mateo, 6 undiagnosed neuro- developmental syndrome	Phenotype timeline, EEG, MRI, family history, genome sequencing, prior negative workup	Phenotype-variant ranking and literature synthesis narrow candidate mechanisms and guide targeted confirmatory testing.
Leila, 34 Crohn's disease on biologics	Symptoms, calprotectin, medication adherence, wearables, sleep, diet, clinician notes	Within-person pattern detection turns fragmented self-tracking into actionable flare prediction and better-timed escalation.

Figure 2: Concrete personas for a computational health stack. The same design principle appears across specialties: assemble the full timeline, compute over it repeatedly, and use the result to improve the next decision.

3. **Chronic inflammatory pathway.** A patient with a relapsing condition integrates symptoms, labs, wearable data, and medication exposure over time. Compute is used within-person, not only across populations, to identify recurrent precursors of flare, intolerance, or non-response.

8 Evaluation and endpoints

If the computational view is correct, it should change how health AI is evaluated. Average discrimination metrics and point-estimate accuracy are not enough. In high-stakes medicine, calibration matters as much as accuracy. Measuring how well a system quantifies its own uncertainty can be more valuable than merely evaluating its top-choice correctness: a system that outputs "60% diagnosis A, 25% diagnosis B, and here is the test that discriminates between them" is radically more useful than a bare point estimate of A, even at identical baseline accuracy. This connects directly to compute elasticity. Sometimes more reasoning does not change the top-ranked answer but dramatically sharpens the uncertainty estimate around it, which clinically determines whether a clinician will order a confirmatory test, refer to a specialist, or treat empirically. The computational view must therefore care about decision-relevant uncertainty at least as much as it cares about the prediction itself. In the domains highlighted here, the relevant question is whether synthesis quality—including calibration—improved at the point where the case could still change direction.

Study type	What changed	Why it matters for the thesis
Pragmatic randomized clinical trial in cardiovascular screening	AI-enabled ECG alerts were tested as an intervention rather than a retrospective classifier [21]	Supports the claim that computational systems can change downstream decisions and outcomes, not merely generate scores.
Pragmatic cluster-randomized deterioration trial	Real-time surveillance for patient deterioration was tested across live inpatient settings [23]	Strengthens the case that high-context clinical surveillance can function as an operational care system.
Patient-record sharing study	Patients with transparent notes shared them with family and caregivers, extending the practical reach of the record [16]	Shows that once patients have access to records, those records begin acting as coordination infrastructure.
Patient-side LLM proof of concept	Patients used large language models on open notes to better understand and work with their records [28]	Gives direct support to the argument that personal compute can become part of the patient toolkit.

Table 3: Evidence that the computational view extends beyond retrospective prediction.

Illustrative translational case

A widely discussed 2025 media report described how an Australian entrepreneur used sequencing, variant analysis, and AI-assisted reasoning to help design a personalized mRNA treatment strategy for his dog Rosie, who had a mast cell tumour [1]. This is not high-level evidence in the way a randomized trial is. It is useful for a different reason: it makes visible what the computational workflow looks like once motivated people can assemble data, models, and synthesis outside a traditional institutional pipeline.

9 What should be built

The implication is not that every patient should be left alone with a chatbot. The implication is that health systems, startups, and regulators should prioritize tools that make high-burden reasoning portable, inspectable, and patient-aligned.

1. **Patient-assembled case files.** Systems should ingest records from multiple providers and preserve chronology rather than flattening everything into isolated visits.
2. **Explicit uncertainty handling.** Models should separate evidence, inference, and missing information, especially in oncology, genetics, psychiatry, and infection.
3. **Decision-oriented outputs.** The output should not just summarize. It should rank the next tests, consultations, or therapeutic branches.
4. **Re-runnable reasoning.** Every major disease course changes over time. The right unit is a living case model, not a one-time answer.
5. **Patient-clinician collaboration.** The design target is a stronger clinical conversation, not consumer isolation.

Domain	Primary endpoint candidates	Why these endpoints fit the computational thesis
Rare disease	Time to diagnosis; diagnostic yield; avoided procedures	Measures whether synthesis reduced time spent on the wrong branch of care.
Precision oncology	Time to molecular recommendation; trial-match yield; change in treatment ranking; time to next line	Measures whether computation improved branching decisions rather than documentation alone.
Sepsis / ICU	Time to escalation; antibiotics or intervention timing; false alerts per patient-day; mortality or ICU transfer	Measures whether longitudinal synthesis improved timing without overwhelming workflow.
Chronic disease / N-of-1 care	Time to flare detection; within-person treatment response; symptom burden; adherence to preferred regimen	Measures whether repeated computation improves individualized management over time.
Patient-facing record synthesis	Patient comprehension; visit preparedness; question quality; fewer duplicated tests; continuity across sites	Measures whether access plus reasoning improved the practical utility of records.

Table 4: Evaluation endpoints aligned with a computational view of medicine. The goal is to measure whether reasoning improves decisions, not only whether predictions look accurate offline.

6. **N-of-1 evidence layers.** Systems should support repeated within-person hypothesis testing when disease is relapsing, heterogeneous, or highly treatment-sensitive.

10 Risks and constraints

The computational view does not weaken the need for evidence or clinical oversight. It sharpens it. High-burden disease is precisely where hallucinated synthesis, hidden bias, dataset shift, and brittle prompting are most dangerous. Better reasoning systems must therefore be auditable, source-grounded, and explicit about what they do not know.

There is also a distributional risk: if compute meaningfully improves outcomes, unequal access to compute will become a health inequity. Data access without reasoning access will be incomplete empowerment. A future health stack that gives patients records but not useful synthesis will still leave too much leverage with the institution.

There is also a harder limit that the computational view must acknowledge honestly. Some biological processes may be computationally irreducible: no model, however large, can predict their trajectory faster than the process itself unfolds. Immune response to a novel pathogen, idiosyncratic drug reactions in a patient with unusual comorbidities, the nonlinear dynamics of some cancers—these may resist compression into predictive shortcuts regardless of how much inference is applied. The framework developed here is explicit about this. Compute elasticity is not universally high. In domains where the biology is genuinely unpredictable or the action space is narrow, larger inference budgets will not help. The computational view

is a claim about where the bottleneck lies *when it lies in reasoning*, not a claim that reasoning is always the bottleneck.

The paper also argues against a common failure mode in AI health discourse: treating “personalization” as a vague user-experience feature. In the computational paradigm, personalization is a property of the reasoning problem itself. It requires timeline integrity, provenance, repeated re-analysis, and outcome feedback.

11 Operational consequences

If this view is correct, the agenda for health AI changes. Product teams should stop treating every workflow as though it deserves the same level of model complexity. Clinical systems should separate protocol problems from search problems. Regulators should ask not only whether a model is accurate on average, but whether it improves the quality of reasoning in the cases where reasoning is the bottleneck. And patient-facing tools should be built around timeline integrity, evidence provenance, and re-runnable case synthesis rather than one-shot conversational convenience.

The economics reinforce this conclusion. A rare-disease diagnostic odyssey averages seven years and accumulates tens of thousands of dollars in unnecessary specialist visits, repeated tests, and procedures on the wrong branch of care. An AI reasoning system that compresses that timeline to months might cost tens of dollars in compute. In oncology, a missed clinical trial match means a patient who might have responded to an experimental therapy never receives it; the value of a correct match—to the patient, the health system, and the trial—is orders of magnitude larger than the cost of the computation that could have found it. In psychiatry, each failed antidepressant trial costs six to eight weeks of continued illness; pharmacogenomic testing costs roughly two hundred dollars and can eliminate likely failures before they are tried. The ratio between the cost of reasoning and the cost of its absence is not two-to-one. It is routinely one hundred to one or higher. This asymmetry makes the computational paradigm feel not aspirational but economically inevitable.

There is also a structural consequence that goes beyond efficiency. Medicine today is organized around the assumption that reasoning is expensive and scarce—it is done by highly trained humans with limited time. This scarcity shapes the referral hierarchy, the triage system, the fifteen-minute consultation, the division of labor between generalists and specialists, and the requirement that patients physically travel to access a subspecialist’s judgment. If AI makes high-quality reasoning cheap and portable, these structures should change. Referrals become less necessary when a generalist has access to specialist-level synthesis. The consultation becomes less about the doctor racing to integrate data in real time and more about judgment and shared decision-making over pre-computed analysis. The medical record transforms from a document that is stored and retrieved into a computational substrate that is actively reasoned over. And geography becomes less binding: a patient in a rural clinic can access the same depth of synthesis as one at an academic medical center. The binding constraint shifts from proximity to experts to data quality and action capacity. This is the deeper paradigm shift the computational view implies: not merely that AI helps doctors, but that cheap computation restructures how medical reasoning is distributed across people, machines, time, and institutions.

The practical payoff is concentration. More compute should go to the places where it buys something real: fewer missed driver mutations, faster rare-disease diagnoses, better antibiotic choices, earlier sepsis escalation, more systematic medication sequencing in psychiatry, cleaner preparation for specialist visits, and better continuity when care fragments across institutions. Results from pragmatic intervention studies in ECG alerting, patient deterioration surveillance, and electronic symptom monitoring suggest that this shift can move from theory into deployed operations when the system is attached to a real decision path [4, 21, 23].

Testable predictions

Prediction 1. Within five years, the most effective clinical AI deployments will increasingly be evaluated on time-to-decision and branch-quality endpoints rather than stand-alone discrimination metrics.

Prediction 2. The first widely adopted patient-facing AI tools in medicine will succeed not as synthetic clinicians but as timeline integrators and case-preparation systems.

Prediction 3. N-of-1 computational care loops will emerge first in rare disease, oncology, and relapsing chronic disease, where within-person variation and high decision burden are most visible.

Prediction 4. Institutions that expose records without exposing usable reasoning will create a new form of incomplete patient empowerment.

12 Conclusion

Medicine is already a computational process. It has been one since the first clinician held two observations in mind and asked which diagnosis fit better. Evidence-based medicine formalized this at the population level. The computational view proposed here formalizes it at the patient level, and argues that the gap between what is currently computed and what could be computed is the single largest source of avoidable harm in complex disease.

The diseases that matter most under this view are not simply the most prevalent. They are the ones that become search problems when viewed at patient resolution—where the hypothesis space is large, the evidence is distributed, and the quality of reasoning directly determines whether the patient ends up on the right branch of care or the wrong one. Precision oncology, rare-disease genomics, resistant infection, psychiatric treatment selection, selected critical care settings, and chronic relapsing diseases are early proof that this lens is practical rather than rhetorical.

The deeper implication is structural. If reasoning has historically been scarce and expensive—bound to the time, memory, and attention of individual clinicians—then cheap, portable computation changes not only the quality of individual decisions but the architecture of medicine itself: how referrals work, what consultations are for, who can access specialist-level synthesis, and whether the medical record is a static archive or a living computational substrate.

That is why a computational view of medicine matters. It provides a principled way to measure where additional inference can genuinely change outcomes, a framework for directing resources toward the cases where the reasoning deficit is largest, and a vision for how health understanding can become a shared, longitudinal, patient-held capability rather than a privilege of the institution. The question is no longer whether medicine will become more computational. It is how fast, how equitably, and how well.

References

- [1] ABC Health Report. How a dog owner used AI to help design a cancer treatment for rosie. <https://www.abc.net.au/listen/programs/healthreport/dog-owner-ai-cancer-treatment/105462166>, 2025. URL <https://www.abc.net.au/listen/programs/healthreport/dog-owner-ai-cancer-treatment/105462166>. Accessed 2025-09-30.
- [2] Roy Adams, Katharine E. Henry, Anirudh Sridharan, Hossein Soleimani, Andong Zhan, Nishi Rawat, Lauren Johnson, David N. Hager, Sara E. Cosgrove, Andrew Markowski, et al. Prospective, multi-site study of patient outcomes after implementation of the TREWS machine learning-based early warning system for sepsis. *Nature Medicine*, 28(7):1455–1460, 2022. doi: 10.1038/s41591-022-01894-0.

- [3] Emily F. Augustine, Melissa Chiong, Ans de Wilde, Eran Elinav, David C. Fajgenbaum, Terence R. Flotte, David B. Goldstein, Rachelle C. Green, Melissa A. Haendel, Lawrence T. Mazzola, et al. A roadmap for n-of-1 therapies. *Nature Reviews Drug Discovery*, 23(7):507–509, 2024. doi: 10.1038/d41573-024-00087-3.
- [4] Ethan Basch, Amy M. Deal, Mark G. Kris, Howard I. Scher, Clifford A. Hudis, Paulina Sabbatini, Lauren Rogak, Ann V. Bennett, Allison C. Dueck, Timothy M. Atkinson, et al. Symptom monitoring with patient-reported outcomes during routine cancer treatment: A randomized controlled trial. *Journal of Clinical Oncology*, 34(6):557–565, 2016. doi: 10.1200/JCO.2015.63.0830.
- [5] Ethan Basch, Deborah Schrag, Sara Henson, Jan Jansen, Benjamin Ginos, Ann Stover, Paula Carr, Patricia Spears, Christopher Jonas, Melissa Weiser, et al. Effect of electronic symptom monitoring on patient-reported outcomes among patients with metastatic cancer: A randomized clinical trial. *JAMA*, 318(2):197–198, 2017. doi: 10.1001/jama.2017.7156.
- [6] Aaron T. Beck and Emily A. P. Haigh. Advances in cognitive theory and therapy: The generic cognitive model. *Annual Review of Clinical Psychology*, 10:1–24, 2014. doi: 10.1146/annurev-clinpsy-032813-153734.
- [7] Rui Chen, George I. Mias, Jennifer Li-Pook-Than, Long Jiang, Hugo Y. K. Lam, Rui Chen, Eyal Miriami, Konrad J. Karczewski, Muthu Hariharan, Frederick E. Dewey, et al. Personal omics profiling reveals dynamic molecular and medical phenotypes. *Cell*, 148(6):1293–1307, 2012. doi: 10.1016/j.cell.2012.02.009.
- [8] Pat Croskerry. A universal model of diagnostic reasoning. *Academic Medicine*, 84(8):1022–1028, 2009. doi: 10.1097/ACM.0b013e3181ace703.
- [9] Tom Delbanco, Jan Walker, Jonathan D. Darer, Joann G. Elmore, Henry J. Feldman, Suzanne G. Leveille, James D. Ralston, Stephen E. Ross, Elisabeth Vodicka, and Valerie D. Weber. Open notes: Doctors and patients signing on. *Annals of Internal Medicine*, 153(2):121–125, 2010. doi: 10.7326/0003-4819-153-2-201007200-00008.
- [10] Mark L. Graber, Nancy Franklin, and Ruthanna Gordon. Diagnostic error in internal medicine. *Archives of Internal Medicine*, 165(13):1493–1499, 2005. doi: 10.1001/archinte.165.13.1493.
- [11] Gordon Guyatt, David Sackett, Jonathan Adachi, Roger Roberts, Jose Chong, David Rosenbloom, and Joseph Keller. A clinician's guide for conducting randomized trials in individual patients. *Canadian Medical Association Journal*, 134(8):889–892, 1986.
- [12] Iva Hassan, Alastair van Heerden, Bala Padmanabhan, Norbert C. Benda, Emanuele Mazzone, Leo Anthony Celi, and Claudia C. Dobler. How can artificial intelligence support shared decision making? a systematic review of ai-enabled decision aids. *npj Digital Medicine*, 7(1):46, 2024. doi: 10.1038/s41746-024-01056-1.
- [13] J. Kevin Hicks, Jeffrey R. Bishop, Katrin Sangkuhl, Daniel J. Müller, Yuan Ji, Susan G. Leckband, J. Steven Leeder, Rochelle L. Graham, David L. Chiuber, Adrián LLerena, Todd C. Skaar, Stuart A. Scott, Julia C. Stingl, Teri E. Klein, Kelly E. Caudle, and Andrea Gaedigk. Clinical Pharmacogenetics Implementation Consortium (CPIC) guideline for CYP2D6 and CYP2C19 genotypes and dosing of selective serotonin reuptake inhibitors. *Clinical Pharmacology & Therapeutics*, 98(2):127–134, 2015. doi: 10.1002/cpt.147.
- [14] Kit Huckvale, Seema Venkatesh, and Helen Christensen. Toward clinical digital phenotyping: a timely opportunity to consider purpose, quality, and safety. *npj Digital Medicine*, 2(1):88, 2019. doi: 10.1038/s41746-019-0166-1.

- [15] Greg Irving, Ana Luisa Neves, Hajira Dambha-Miller, Ai Oishi, Hiroko Tagashira, Anistassia Verho, and John Holden. International variations in primary care physician consultation time: a systematic review of 67 countries. *BMJ Open*, 7(10):e017902, 2017. doi: 10.1136/bmjopen-2017-017902.
- [16] Sara L. Jackson, Roanne Mejilla, Jonathan D. Darer, Natalia V. Oster, James D. Ralston, Suzanne G. Leveille, Jan Walker, and Tom Delbanco. Patients who share transparent visit notes with others: characteristics, risks, and benefits. *Journal of Medical Internet Research*, 16(11):e247, 2014. doi: 10.2196/jmir.3363.
- [17] Heather S. L. Jim, Andrea I. Hoogland, Noah C. Brownstein, Andre Barata, Adam P. Dicker, Hans Knoop, Brenda D. Gonzalez, Robert Perkins, Dana Rollison, Rita Nanda, et al. Innovations in research and clinical care using patient-generated health data. *CA: A Cancer Journal for Clinicians*, 70(3): 182–199, 2020. doi: 10.3322/caac.21608.
- [18] Aafke H. Jonker, Costas Demetzos, Anthonius de Boer, Eric G. E. de Vries, Johan T. den Dunnen, Marc T. W. Handley, Anke W. Hoes, Hubert G. M. Leufkens, Jean J. Sikkens, Rosalie van der Graaf, et al. N-of-1 medicine and treatment outcomes: the future is bright. *Nature Reviews Drug Discovery*, 24(5):335–336, 2025. doi: 10.1038/d41573-025-00004-z.
- [19] John Jumper, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger, Kathryn Tunyasuvunakool, Russ Bates, Augustin Žídek, Anna Potapenko, et al. Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873):583–589, 2021. doi: 10.1038/s41586-021-03819-2.
- [20] Elizabeth O. Lillie, Bence Patay, Judith Diamant, Brian Issell, Eric J. Topol, and Nicholas J. Schork. The n-of-1 clinical trial: the ultimate strategy for individualizing medicine? *Personalized Medicine*, 8(2):161–173, 2011. doi: 10.2217/pme.11.7.
- [21] Chin-Sheng Lin, Wei-Ting Liu, Dung-Jang Tsai, Yu-Sheng Lou, Chiao-Hsiang Chang, Chiao-Chin Lee, Wen-Hui Fang, Chih-Chia Wang, Yen-Yuan Chen, Wei-Shiang Lin, et al. Ai-enabled electrocardiography alert intervention and all-cause mortality: a pragmatic randomized clinical trial. *Nature Medicine*, 30(5):1461–1470, 2024. doi: 10.1038/s41591-024-02961-4.
- [22] Birger Moëll, Fredrik Sand Aronsson, and Sanian Akbar. Medical reasoning in LLMs: an in-depth analysis of DeepSeek R1. *Frontiers in Artificial Intelligence*, 8:1616145, 2025. doi: 10.3389/frai.2025.1616145.
- [23] Rebecca Moses, Lu Chen, Irtiza Khan, Zachary Pruitt, Andrew Sharp, Yu Ruan, Yiyun He, Jason Stifter, Hossein Estiri, Stephen A. Collins, et al. Real-time surveillance system for patient deterioration: a pragmatic cluster-randomized controlled trial. *Nature Medicine*, 2025. doi: 10.1038/s41591-025-03609-7.
- [24] Linda G. Park, Felicity Ng, Jimmy Shim, Adham Elnaggar, Oscar Villero, Frank Perna, Steven M. Paul, and Christopher R. Friese. Perceptions and experiences of using mobile technology for symptom management in patients with cancer: a systematic review. *JMIR mHealth and uHealth*, 6(1):e30, 2018. doi: 10.2196/mhealth.8748.
- [25] Brian D. Piening, Wenyu Zhou, Kevin Contrepois, Hannes Röst, Gabriela J. Gu Urban, Tanzeem Mishra, Brett M. Hanson, Elizabeth J. Bautista, Steven R. Leopold, Chih-Ying Yeh, et al. Integrative personal omics profiles during periods of weight gain and loss. *Cell Systems*, 6(2):157–170.e8, 2018. doi: 10.1016/j.cels.2018.01.013.
- [26] Alvin Rajkomar, Eyal Oren, Kai Chen, Andrew M. Dai, Nissan Hajaj, Michaela Hardt, Peter J. Liu, Xiaobing Liu, Jake Marcus, Mimi Sun, et al. Scalable and accurate deep learning with electronic health records. *npj Digital Medicine*, 1(1):18, 2018. doi: 10.1038/s41746-018-0029-1.

- [27] Ugur Sahin, Evelyn Derhovanessian, Michael Miller, Ben Kloke, Peter Simon, Markus Löwer, Thomas Bukur, Anat Tadmor, Ulrike Luxemburger, Benedikt Schrörs, et al. Personalized rna mutanome vaccines mobilize poly-specific therapeutic immunity against cancer. *Nature*, 547(7662):222–226, 2017. doi: 10.1038/nature23003.
- [28] Liz Salmi, Dana M. Lewis, Jennifer L. Clarke, Zhiyong Dong, Rudy Fischmann, Emily I. McIntosh, Chethan R. Sarabu, and Catherine M. DesRoches. A proof-of-concept study for patient use of open notes with large language models. *JAMIA Open*, 8(2), 2025. doi: 10.1093/jamiaopen/ooaf021.
- [29] Nicholas J. Schork. Personalized medicine: Time for one-person trials. *Nature*, 520(7549):609–611, 2015. doi: 10.1038/520609a.
- [30] Jonathan M. Stokes, Kevin Yang, Kyle Swanson, Wengong Jin, Andrés Cubillos-Ruiz, Nicole M. Donghia, Craig R. MacNair, Shawn French, Lindsey A. Carfrae, Zoe Bloom-Ackermann, et al. A deep learning approach to antibiotic discovery. *Cell*, 180(4):688–702.e13, 2020. doi: 10.1016/j.cell.2020.01.021.
- [31] Eric J. Topol. High-performance medicine: the convergence of human and artificial intelligence. *Nature Medicine*, 25(1):44–56, 2019. doi: 10.1038/s41591-018-0300-7.
- [32] Katherine R. Tuttle, Radica Z. Alicic, O. Kenrik Duru, C. Raymond Jones, Karen B. Daratha, Stephanie B. Nicholas, and Joshua J. Neumiller. Clinical characteristics of and risk factors for chronic kidney disease among adults and perspectives on patient-centered precision medicine. *Kidney International*, 99(1):52–66, 2021. doi: 10.1016/j.kint.2020.10.046.
- [33] Monica H. Wojcik, Gabrielle Lemire, Eva Berger, Maha S. Zaki, Mariel Wissmann, Wathone Win, Susan M. White, Ben Weisburd, Dagmar Wieczorek, Leigh B. Waddell, et al. Genome sequencing for diagnosing rare diseases. *The New England Journal of Medicine*, 390(21):1985–1997, 2024. doi: 10.1056/NEJMoA2314761.